


Extreme Gradient Boosting

  
XGBoost




*XGboost is simpler than Deep Learning so before trying Deep Learning, it's worth trying XGboost.*

# Tree-Based Models

*Ensemble modeling*


single



- cut points  
- order of features


1 iteration

bagging



parallel

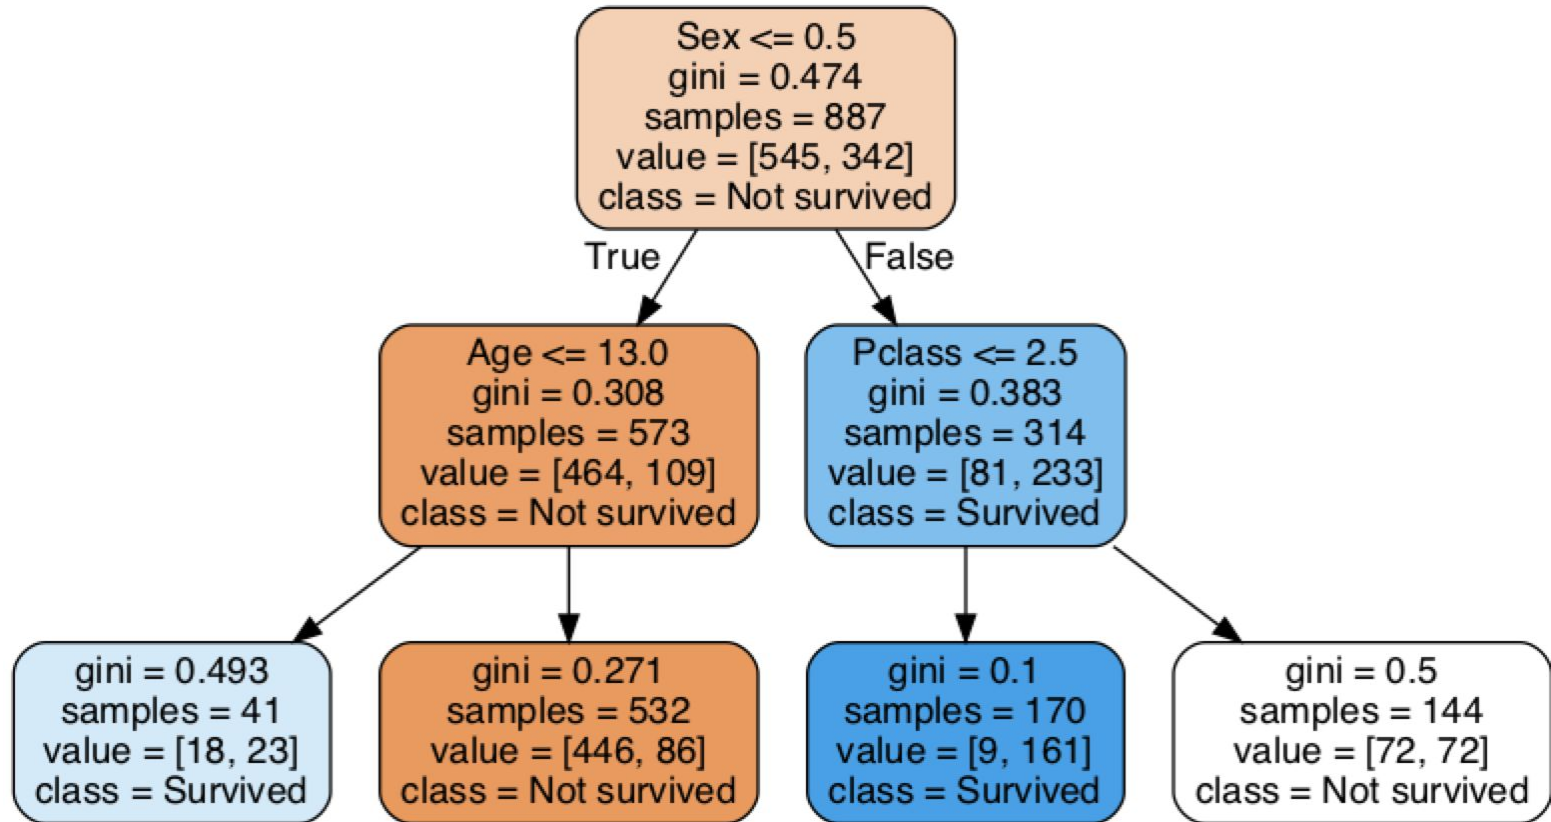
boosting



sequential

*Result is average of each bag.  
Since it's parallel, it's faster than boosting.*

# Single Decision Tree




# Bagging vs Boosting

Bagging algorithms:

- Bagging Ensemble
- Random Forest
- Extra Trees

Boosting algorithms:


- Adaboost
- Gradient Boosting
- XGBoost



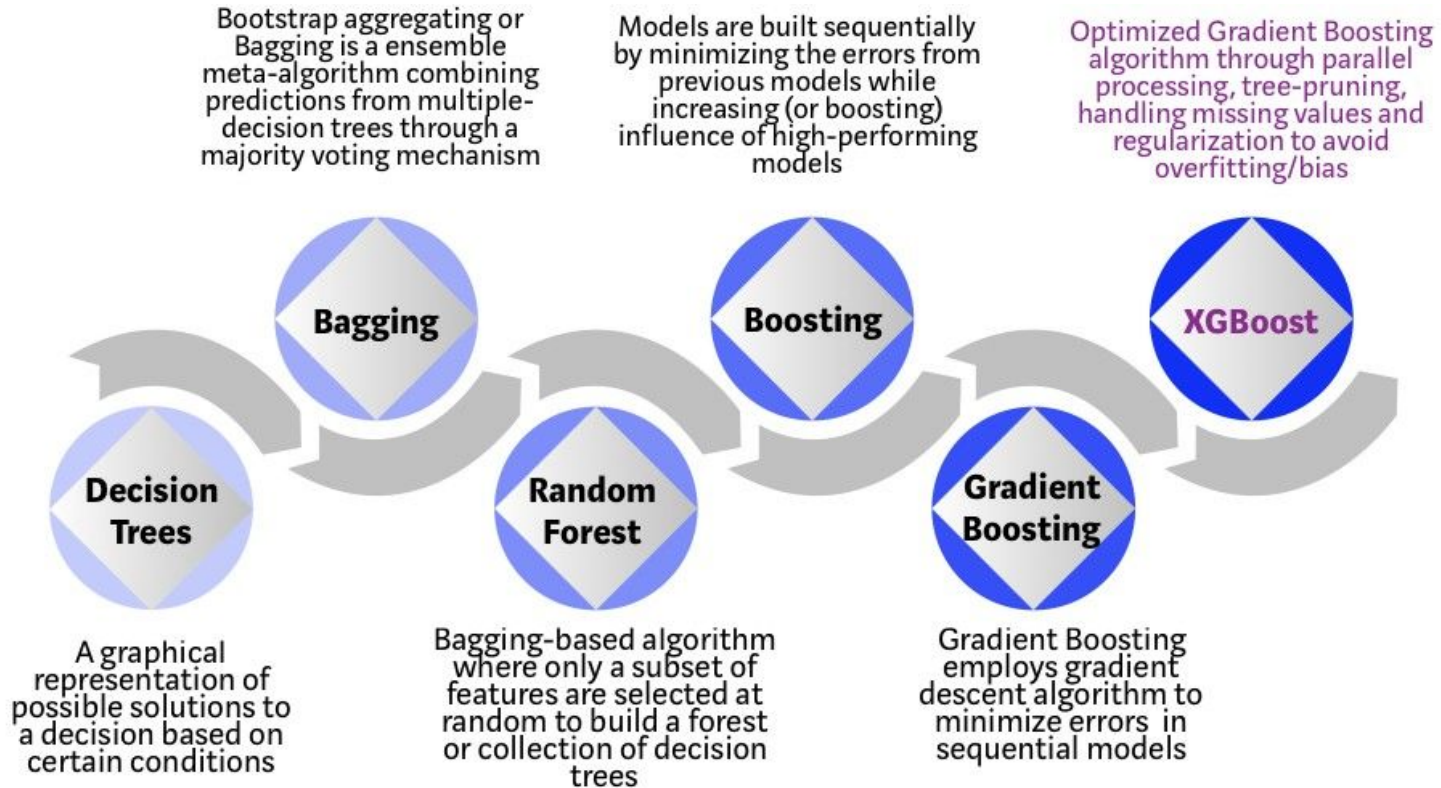
# Random Forest

- Bootstrap samples
- Random selection of features
- Run in parallel

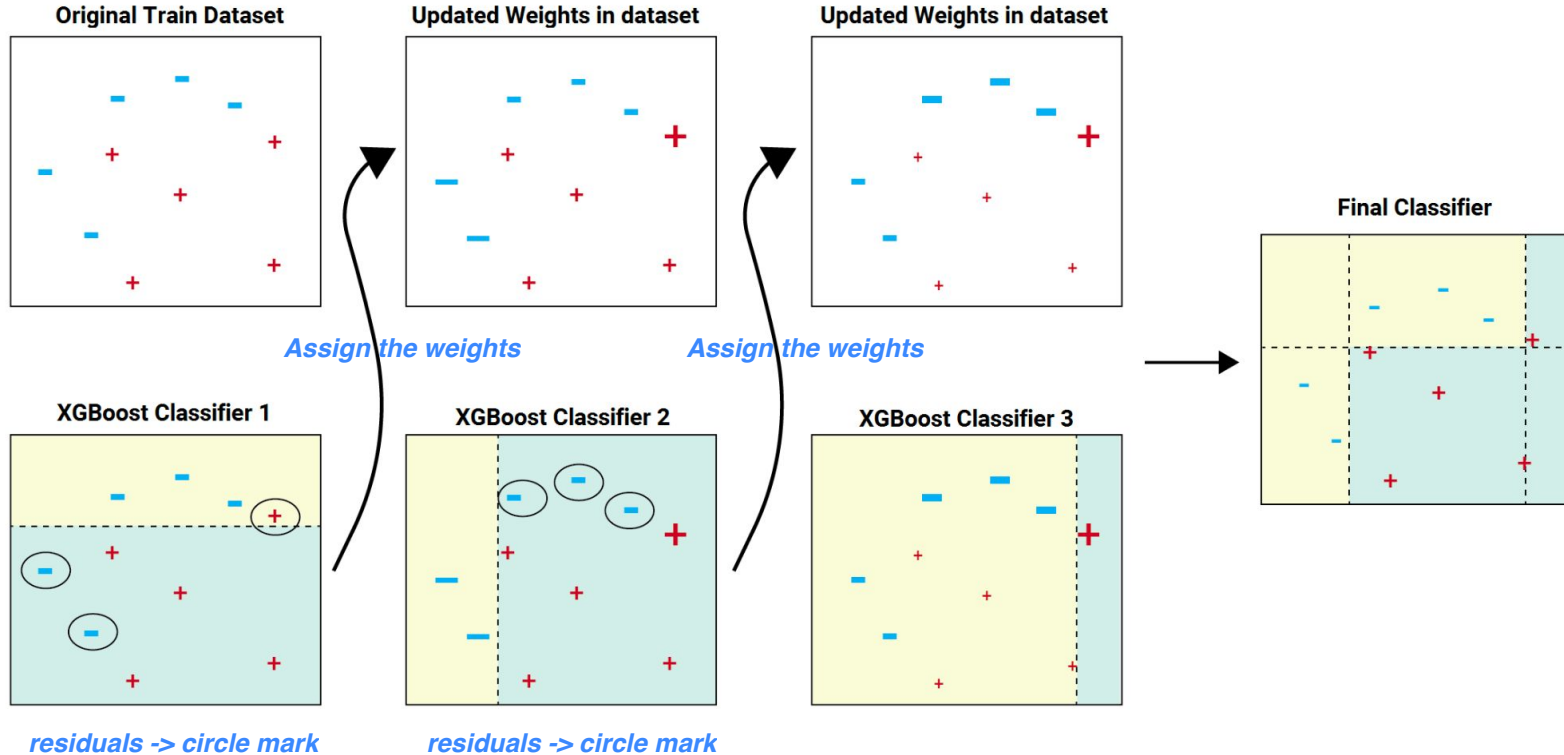
- *increases generability*  
- *decreases overfit*



# Bagging vs Boosting

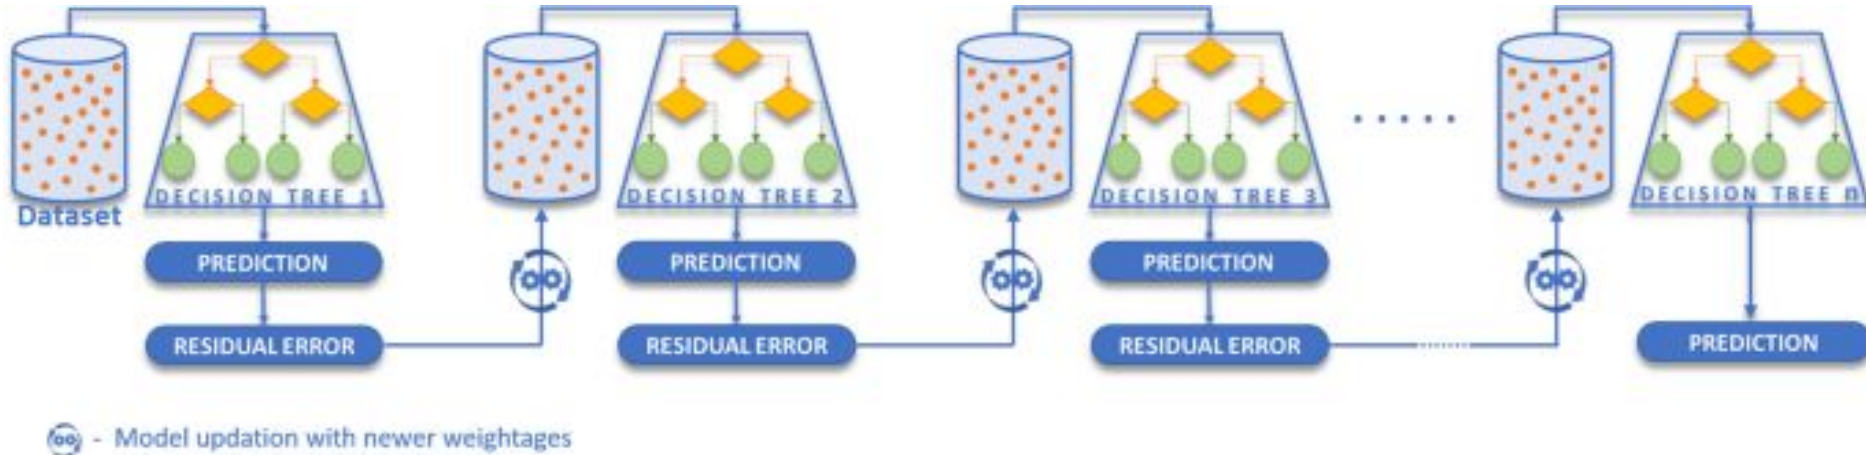


# XGBoost (Extreme Gradient Boosting)





# XGBoost models the residuals





# Hyperparameter tuning: Random Forest

```
rf_model = RandomForestClassifier(max_depth=8,  
                                min_samples_leaf=10,  
                                n_estimators=100,  
                                max_features='sqrt' # max_features=sqrt(n_features)|  
                                criterion='gini'  
                                )
```

- max\_depth: (default none): The maximum depth of the tree.
- min\_samples\_leaf (default=1): The minimum number of samples required to be at a leaf node.
- n\_estimators (default=100): The number of trees in the forest.
- max\_features {"sqrt", "log2", None}: The number of features to consider when looking for the best split
- n\_jobs (default: none): The number of jobs to run in parallel
- class\_weight{"balanced", "balanced\_subsample"}, default=None

# Hyperparameter tuning: XGBoost

```
xgb_model = XGBClassifier(max_depth=6,  
                           min_child_weight=1,  
                           gamma=0,  
                           subsample=1,  
                           learning_rate=0.3)
```

- **min\_child\_weight** (default=1): Minimum sum of instance weight (hessian) needed in a child.
- **gamma** (default=0): Minimum loss reduction required to make a further partition on a leaf node of the tree. The larger gamma is, the more conservative the algorithm will be.
- **subsample** (default=1): Subsample ratio of the training instances. Setting it to 0.5 means that XGBoost would randomly sample half of the training data prior to growing trees
- **learning\_rate** (default=0.3): Step size shrinkage used in update to prevents overfitting.

# Two ways to control overfitting in XGBoost

1. Directly control model complexity.
  - This includes `max_depth`, `min_child_weight` and `gamma`.
2. Add randomness to make training robust to noise.
  - This includes `subsample` and `colsample_bytree`.
  - You can also reduce stepsize `eta`. Remember to increase `num_round` when you do so.

[https://xgboost.readthedocs.io/en/stable/tutorials/param\\_tuning.html](https://xgboost.readthedocs.io/en/stable/tutorials/param_tuning.html)